



Project Title	Twitter Financial News
Tools	Visual Studio code / jupyter notebook
Domain	Finance Analyst
Project Difficulties level	intermediate

Dataset : Dataset is available in the given link. You can download it at your convenience.

[Click here to download data set](#)

About the Data

The Twitter Financial News dataset is an English-language dataset containing an annotated corpus of finance-related tweets. This dataset is used to classify finance-related tweets for their topic.

Featured Notebooks:

Ideas:

1. The data is a multi-label text classification problem with imbalanced data.
2. The links within the texts could be extracted.

3. EDA

The dataset holds 21,107 documents annotated with 20 labels:

"LABEL_0": "Analyst Update",
"LABEL_1": "Fed | Central Banks",
"LABEL_2": "Company | Product News",
"LABEL_3": "Treasuries | Corporate Debt",
"LABEL_4": "Dividend",
"LABEL_5": "Earnings",
"LABEL_6": "Energy | Oil",
"LABEL_7": "Financials",
"LABEL_8": "Currencies",
"LABEL_9": "General News | Opinion",
"LABEL_10": "Gold | Metals | Materials",
"LABEL_11": "IPO",
"LABEL_12": "Legal | Regulation",
"LABEL_13": "M&A | Investments",
"LABEL_14": "Macro",
"LABEL_15": "Markets",

"LABEL_16": "Politics",

"LABEL_17": "Personnel Change",

"LABEL_18": "Stock Commentary",

"LABEL_19": "Stock Movement"

Example: You can get the basic idea how you can create a project from here

Creating a **Twitter Financial News Analysis** project involves building a pipeline that processes, analyzes, and models financial sentiments. Here's a step-by-step explanation tailored for your experience level.

Step 1: Problem Definition

- **Objective:** Analyze Twitter financial news and classify sentiments (e.g., **Positive**, **Neutral**, **Negative**) or financial impacts (e.g., **Bullish**, **Bearish**).
 - **Data:**
 - **text:** Tweets or news text related to financial markets.
 - **label:** The sentiment or impact associated with each text (categorical values).
-

Step 2: Data Collection

Assume we have a dataset, or use the **Tweepy API** to collect financial tweets.

Code for Collecting Data:

```
python
code
import tweepy
import pandas as pd

# Twitter API credentials
```

```
api_key = 'your_api_key'
api_secret = 'your_api_secret'
access_token = 'your_access_token'
access_token_secret = 'your_access_token_secret'

# Authenticate
auth = tweepy.OAuth1UserHandler(api_key, api_secret,
access_token, access_token_secret)
api = tweepy.API(auth)

# Fetch tweets
query = 'stock market' # Financial keyword
tweets = tweepy.Cursor(api.search_tweets, q=query, lang="en",
tweet_mode='extended').items(1000)

# Store tweets in DataFrame
data = {'text': [], 'label': []} # Add labels manually or use
pre-annotated data
for tweet in tweets:
    data['text'].append(tweet.full_text)

df = pd.DataFrame(data)
df.to_csv('financial_tweets.csv', index=False)
```

Step 3: Data Cleaning

Clean the tweets to remove irrelevant information like URLs, hashtags, and mentions.

Code:

python

code

```
import re

def clean_text(text):
    text = re.sub(r'http\S+|www\S+|https\S+', '', text,
flags=re.MULTILINE) # Remove URLs
    text = re.sub(r'\@\w+|\#', '', text) # Remove mentions and
hashtags
    text = re.sub(r'^\w\s]', '', text) # Remove punctuation
    text = re.sub(r'\d+', '', text) # Remove numbers
    text = text.lower() # Convert to lowercase
    return text

df['clean_text'] = df['text'].apply(clean_text)
print(df.head())
```

Step 4: Exploratory Data Analysis (EDA)

4.1 Sentiment Distribution

Visualize the distribution of sentiment labels.

Code:

python

code

```
import matplotlib.pyplot as plt
import seaborn as sns

# Sentiment distribution
sns.countplot(data=df, x='label', palette='viridis')
plt.title('Sentiment Distribution')
plt.xlabel('Sentiment')
plt.ylabel('Count')
plt.show()
```

4.2 Word Cloud

Visualize frequent words in tweets.

python

code

```
from wordcloud import WordCloud

# Generate Word Cloud
text = ' '.join(df['clean_text'])
wordcloud = WordCloud(width=800, height=400,
background_color='white').generate(text)
```

```
# Display Word Cloud
plt.figure(figsize=(10, 5))
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.show()
```

Step 5: Feature Engineering

5.1 Text Vectorization

Use **TF-IDF** to convert text into numerical features.

Code:

python
code

```
from sklearn.feature_extraction.text import TfidfVectorizer

# TF-IDF Vectorization
vectorizer = TfidfVectorizer(max_features=5000) # Top 5000
words
X = vectorizer.fit_transform(df['clean_text']).toarray()
y = df['label'] # Target variable
```

5.2 Split Data

python

code

```
from sklearn.model_selection import train_test_split

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

Step 6: Model Training

6.1 Use Logistic Regression

Logistic Regression is a good starting point for classification problems.

Code:

python

code

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report,
confusion_matrix

# Train the model
model = LogisticRegression()
model.fit(X_train, y_train)

# Predict on test data
y_pred = model.predict(X_test)
```

```
# Evaluate the model
print(classification_report(y_test, y_pred))
```

Step 7: Data Visualization

7.1 Confusion Matrix

python

code

```
import seaborn as sns
import matplotlib.pyplot as plt

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=model.classes_, yticklabels=model.classes_)
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```

Step 8: Advanced Modeling

Experiment with more sophisticated models like **BERT** or **LSTM**.

Using Hugging Face's BERT:

python

code

```
from transformers import BertTokenizer,
BertForSequenceClassification
from torch.utils.data import DataLoader, Dataset

# Tokenize text using BERT tokenizer
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
tokens = tokenizer(list(df['clean_text']), padding=True,
truncation=True, return_tensors="pt")

# Use BERT model for classification
model =
BertForSequenceClassification.from_pretrained('bert-base-uncase
d', num_labels=len(df['label'].unique()))
```

Step 9: Deploy the Model

Deploy using **Flask** or **Streamlit** for real-time predictions.

Streamlit Example:

python

code

```
import streamlit as st
```

```
st.title('Twitter Financial News Sentiment Analysis')
user_input = st.text_input("Enter a financial tweet:")
if st.button("Predict"):
    clean_input = clean_text(user_input)
    input_vector =
vectorizer.transform([clean_input]).toarray()
    prediction = model.predict(input_vector)
    st.write(f"Predicted Sentiment: {prediction[0]}")
```

Step 10: Statistical Insights

Analyze relationships between sentiment and financial trends using external data like stock prices.

Example: You can get the basic idea how you can create a project from here

Sample code with output

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from nltk.corpus import stopwords
from nltk import WordNetLemmatizer
from nltk import word_tokenize
import nltk
from wordcloud import WordCloud

import re
import tensorflow as tf
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import
pad_sequences
from tensorflow.keras.layers import Input, Dense, Embedding,
LSTM, Bidirectional, GRU, GlobalMaxPooling1D, Dropout,
SimpleRNN, Conv1D
from tensorflow.keras.models import Model
from tensorflow.keras.optimizers import SGD, Adam
from tensorflow.keras.losses import
```

```
SparseCategoricalCrossentropy, categorical_crossentropy
```

```
from sklearn.feature_extraction.text import CountVectorizer
```

1. Preprocessing

In [46]:

```
# nltk.download("stopwords")
```

```
# nltk.download("punkt")
```

```
# nltk.download("wordnet")
```

In [47]:

```
# get stopwords
```

```
stops = set(stopwords.words("english"))
```

In [48]:

```
df_train =
```

```
pd.read_csv("/kaggle/input/twitter-financial-news/train_data.csv")
```

```
df_test =
```

```
pd.read_csv("/kaggle/input/twitter-financial-news/valid_data.csv")
```

```
In [49]:  
df_train.head()
```

Out[49]:

	text	label
0	Here are Thursday's biggest analyst calls: App...	0
1	Buy Las Vegas Sands as travel to Singapore bui...	0
2	Piper Sandler downgrades DocuSign to sell, cit...	0
3	Analysts react to Tesla's latest earnings, bre...	0

4	Netflix and its peers are set for a 'return to...	0
---	---	---

In [50]:

```
train_labels = df_train["label"]  
train_corpus = df_train["text"]
```

```
test_labels = df_test["label"]  
test_corpus = df_test["text"]
```

In [51]:

get correct mapping of ordinal encoded target variables

```
label_mapping = {"LABEL_0": "Analyst Update",  
  
                 "LABEL_1": "Fed | Central Banks",  
  
                 "LABEL_2": "Company | Product News",  
  
                 "LABEL_3": "Treasuries | Corporate Debt",  
  
                 "LABEL_4": "Dividend",
```


"LABEL_5": "Earnings",

"LABEL_6": "Energy | Oil",

"LABEL_7": "Financials",

"LABEL_8": "Currencies",

"LABEL_9": "General News | Opinion",

"LABEL_10": "Gold | Metals | Materials",

"LABEL_11": "IPO",

"LABEL_12": "Legal | Regulation",

"LABEL_13": "M&A | Investments",

"LABEL_14": "Macro",

"LABEL_15": "Markets",

"LABEL_16": "Politics",

"LABEL_17": "Personnel Change",

```
"LABEL_18": "Stock Commentary",  
  
"LABEL_19": "Stock Movement"  
}
```

In [52]:

```
label_mapping = {k:v for k,v in zip(range(20),  
label_mapping.values())}
```

In [53]:

```
label_mapping
```

Out[53]:

```
{0: 'Analyst Update',  
 1: 'Fed | Central Banks',  
 2: 'Company | Product News',  
 3: 'Treasuries | Corporate Debt',  
 4: 'Dividend',  
 5: 'Earnings',  
 6: 'Energy | Oil',  
 7: 'Financials',
```

```
8: 'Currencies',
9: 'General News | Opinion',
10: 'Gold | Metals | Materials',
11: 'IPO',
12: 'Legal | Regulation',
13: 'M&A | Investments',
14: 'Macro',
15: 'Markets',
16: 'Politics',
17: 'Personnel Change',
18: 'Stock Commentary',

19: 'Stock Movement'}
```

In [54]:

```
# get understanding of the text data
```

```
# number of the random articles
```

```
S = 5
```

```
inds = np.random.choice(train_corpus.index, 5)
```

```
for i in inds:
```

```
    print(train_corpus[i])
```

Klarna's valuation slumps to \$6.7 billion with \$800 million

raise <https://t.co/lrNWSfKGgF> <https://t.co/G0tE1Sv2n9>
WATCH: SAS said that a pilot strike could keep the company from
accessing long-term capital it needs for reorganization and
thus the airline could collapse <https://t.co/m08vgf3CSJ>
<https://t.co/7Y4YMZkXSc>
SILVERSTONE MASTER UK Regulatory Announcement: FRN Variable
Rate Fix <https://t.co/a5eW1kIo89> <https://t.co/K8UEJVjoVX>
Consumer Cable and Internet Spending Plateaus Despite
Cross-category Surge in U.S. Household Expenses
<https://t.co/AZa0TUJ3a6> <https://t.co/eMdIL0xjSU>
RBN Energy's Refined Fuels Analytics Team Welcomes John Auers
<https://t.co/mbbCkBR4fG> <https://t.co/XA1isVTb3Y>

In [55]:

*# there is a need to remove all hyperlinks, since they do not
contain any contextual text data*

```
def remove_hyperlinks_and_punctuation(text):  
    pattern = r'\bhttps?:\/\/\/\S+|[\^\w\s]'  
    new_text = re.sub(pattern, "", text)  
  
    return new_text
```

In [56]:

```
train_corpus_cleaned = train_corpus.apply(lambda x:
remove_hyperlinks_and_punctuation(x))
test_corpus_cleaned = test_corpus.apply(lambda x:
remove_hyperlinks_and_punctuation(x))
```

In [57]:

```
for i in inds:
    print(train_corpus_cleaned[i])
```

Klarnas valuation slumps to 67 billion with 800 million raise
WATCH SAS said that a pilot strike could keep the company from
accessing longterm capital it needs for reorganization and thus
the airline could collapse

SILVERSTONE MASTER UK Regulatory Announcement FRN Variable Rate
Fix

Consumer Cable and Internet Spending Plateaus Despite

Crosscategory Surge in US Household Expenses

RBN Energys Refined Fuels Analytics Team Welcomes John Auers

2. Tokenize corpus

In [58]:

use count vectorizer to feed in ANN without paying attention to the sequence

Could make sense, because classification of the source of the news may not be dependent on the sequence of the words, but rather on some certain important words which are unique for the source

```
class LemmaTokenizer(object):  
    def __init__(self):  
        self.wnl = WordNetLemmatizer()  
    def __call__(self, articles):  
        return [self.wnl.lemmatize(t) for t in  
word_tokenize(articles)]
```

```
MAX_VOCAB_SIZE_ANN = 18000
```

```
vectorizer = CountVectorizer(tokenizer=LemmaTokenizer(),  
stop_words=list(stops), max_features=MAX_VOCAB_SIZE_ANN)  
vectorizer.fit(train_corpus_cleaned)
```

```
train_data_ann = vectorizer.transform(train_corpus_cleaned)  
test_data_ann = vectorizer.transform(test_corpus_cleaned)
```

/opt/conda/lib/python3.7/site-packages/sklearn/feature_extracti

```
on/text.py:517: UserWarning: The parameter 'token_pattern' will
not be used since 'tokenizer' is not None'
```

```
"The parameter 'token_pattern' will not be used"
```

```
/opt/conda/lib/python3.7/site-packages/sklearn/feature_extracti
```

```
on/text.py:401: UserWarning: Your stop_words may be
inconsistent with your preprocessing. Tokenizing the stop words
generated tokens ['d', 'll', 're', 's', 've', 'could',
'doe', 'ha', 'might', 'must', 'n't', 'need', 'sha', 'wa', 'wo',
'would'] not in stop_words.
```

```
% sorted(inconsistent)
```

```
In [59]:
```

```
# create batch generator to feed sparse matrix into keras fit
method
```

```
def batch_generator(X, y, batch_size):
    number_of_batches = X.shape[0]/batch_size
    counter=0
    shuffle_index = np.arange(np.shape(y)[0])
    np.random.shuffle(shuffle_index)
    X = X[shuffle_index, :]
    y = y[shuffle_index]
    while 1:
        # each batch contains all the shuffled indices
```

```

        index_batch = shuffle_index[batch_size*counter :
batch_size*(counter+1)]
        X_batch = X[index_batch,:].todense()
        y_batch = y[index_batch]
        counter += 1
    yield(np.array(X_batch),y_batch)
    if (counter < number_of_batches):
        np.random.shuffle(shuffle_index)
        counter=0

```

In [60]:

Apply lemmatization to corpus - remove maybe later in case of poor performance

used for sequence models

```
wnl = WordNetLemmatizer()
```

```
def lemmatize_sentence(sentence):
```

```
    tokens = [wnl.lemmatize(t) for t in word_tokenize(sentence)]
```

```
    return (" ").join(tokens)
```

```
train_corpus_cleaned = train_corpus_cleaned.apply(lambda x:
lemmatize_sentence(x))
```

```
test_corpus_cleaned = test_corpus_cleaned.apply(lambda x:
lemmatize_sentence(x))
```



```
In [61]:
```

```
for i in inds:
    print(train_corpus_cleaned[i])
```

```
Klarnas valuation slump to 67 billion with 800 million raise
WATCH SAS said that a pilot strike could keep the company from
accessing longterm capital it need for reorganization and thus
the airline could collapse
SILVERSTONE MASTER UK Regulatory Announcement FRN Variable Rate
Fix
Consumer Cable and Internet Spending Plateaus Despite
Crosscategory Surge in US Household Expenses
RBN Energys Refined Fuels Analytics Team Welcomes John Auers
```

```
In [62]:
```

```
# transform the corpus to get training and test data
```

```
MAX_VOCAB_SIZE = 30000
```

```
tokenizer = Tokenizer(num_words=MAX_VOCAB_SIZE,
                       oov_token="OOV",
```

```
lower=True)
```

```
tokenizer.fit_on_texts(train_corpus_cleaned)
```

```
train_data = tokenizer.texts_to_sequences(train_corpus_cleaned)
```

```
test_data = tokenizer.texts_to_sequences(test_corpus_cleaned)
```

```
In [63]:
```

```
# get max len of sentences for padding
```

```
maxlen1 = max(len(sent) for sent in train_data)
```

```
maxlen2 = max(len(sent) for sent in test_data)
```

```
T = max(maxlen1, maxlen2)
```

```
print(T)
```

```
64
```

```
In [64]:
```

```
word2idx = tokenizer.word_index
```

```
V = len(word2idx)
```

```
print("The corpus contains %s words!" % V)
```

The corpus contains 25558 words!

In [65]:

```
# pad sequences
```

```
train_data_padded = pad_sequences(train_data, maxlen=T)
```

```
test_data_padded = pad_sequences(test_data, maxlen=T)
```

In [66]:

```
print("Train tensor shape:", train_data_padded.shape)
```

```
print("Test tensor shape:", test_data_padded.shape)
```

Train tensor shape: (16990, 64)

Test tensor shape: (4117, 64)

In [67]:

```
# get class number
```

```
K = len(set(train_labels))
```

```
print(f"We have %s unique labels!" % K)
```

We have 20 unique labels!

3. Data overview and visualization

In [68]:

```
# show word overview with wordcloud
```

```
# not using all words, but rather subsamples
```

```
random_subsample_word_wordcloud_size =
```

```
round(len(train_corpus_cleaned) * 0.33)
```

```
wordcloud_ids = np.random.choice(len(train_corpus_cleaned),  
random_subsample_word_wordcloud_size)
```

```
concat_text = ""
```


```
for i in wordcloud_ids:
```

```
    concat_text += " " + train_corpus_cleaned[i]
```

```
wordcloud = WordCloud(background_color="white", width=800,  
height=400, stopwords=stops).generate(concat_text)
```

In [69]:

```
fig, ax = plt.subplots(figsize=(14,10))
ax.imshow(wordcloud, interpolation="bilinear")
ax.axis("off")
plt.title("Wordcloud of most frequent words", fontsize=40)
plt.show()
```




In [70]:


distribution of train targets

```
fig, ax = plt.subplots(figsize=(12,8))
sns.countplot(x=train_labels.index, data=train_labels, ax=ax,
```

```
palette="hls")
plt.title("Distribution of train labels", fontsize=25)
plt.show()
print("Highly imbalanced data!!")
```



Distribution of train labels




Highly imbalanced data!!


In [71]:

distribution of train targets

```
fig, ax = plt.subplots(figsize=(12,8))
sns.countplot(x=test_labels.index, data=test_labels, ax=ax,
palette="hls")
plt.title("Distribution of test labels", fontsize=25)
plt.show()
print("Highly imbalanced data!!")
```



Distribution of test labels



Highly imbalanced data!!

In [72]:

```
# highly imbalanced data, so inverse class weights are computed  
for the loss function
```

```
class_weights =
```

```
train_labels.value_counts(normalize=True).sort_index()
```

```
inverse_class_weights = class_weights.apply(lambda x: 1 / x)
```

```
inverse_class_weights
```

Out[72]:

0	66.627451
1	20.298686
2	4.792666
3	52.928349
4	47.325905
5	17.213779
6	32.423664
7	27.227564
8	102.349398
9	10.912010
10	246.231884


```
11      386.136364
12      34.887064
13      36.072187
14       9.324918
15      33.912176
16      17.248731
17      34.323232
18       8.021719
19      20.643985
```

Name: label, dtype: float64

4. Build the model

In [73]:

```
# Embedding dimension
```

```
D = 64
```

```
hidden_states1 = 12
```

```
def make_model_rnn():
```

```
    with tf.device("/GPU:0"):
```

```
        i = Input(shape=(T,))
```

```
        x = Embedding(V + 1, D)(i)
```

```
        x = SimpleRNN(hidden_states1, return_sequences=True)(x)
```

```
        x = GlobalMaxPooling1D()(x)
```

```
        x = Dense(64, activation="relu")(x)
```

```
x = Dropout(0.5)(x)
x = Dense(K, activation="softmax")(x)

model = Model(i,x)
model.compile(loss=SparseCategoricalCrossentropy(),
optimizer=Adam(learning_rate=0.001), metrics=["accuracy"])

return model
```

In [74]:

```
def make_model_cnn():
    with tf.device("/GPU:0"):
        i = Input(shape=(T,))
        x = Embedding(V + 1, D)(i)
        x = Conv1D(32,3)(x)
        x = GlobalMaxPooling1D()(x)
        x = Dense(64, activation="relu")(x)
        x = Dropout(0.3)(x)
        x = Dense(K, activation="softmax")(x)

    model = Model(i,x)
    model.compile(loss=SparseCategoricalCrossentropy(),
optimizer=SGD(learning_rate=0.0001, momentum=0.9),
metrics=["accuracy"])
```

```
return model
```

In [75]:

```
def make_model_ann():  
    with tf.device("/GPU:0"):  
        i = Input(shape=(MAX_VOCAB_SIZE_ANN,))  
        x = Dense(1028, activation="relu")(i)  
        x = Dropout(0.2)(x)  
        x = Dense(512, activation="relu")(x)  
        x = Dropout(0.2)(x)  
        x = Dense(256, activation="relu")(x)  
        x = Dense(K, activation="softmax")(x)  
  
        model = Model(i, x)  
        model.compile(loss=SparseCategoricalCrossentropy(),  
optimizer=SGD(learning_rate=0.0001, momentum=0.9),  
metrics=["accuracy"])  
        # maybe focal loss  
  
    return model
```

5 Fit the model

In [76]:

```
model = make_model_cnn()
```

```
print(model.summary())
```

Model: "model_1"

--		
Layer (type)	Output Shape	Param #
=====		
==		
input_2 (InputLayer)	[(None, 64)]	0
embedding_1 (Embedding)	(None, 64, 64)	1635776
conv1d_1 (Conv1D)	(None, 62, 32)	6176
global_max_pooling1d_1 (GlobalMaxPooling1D)	(None, 32)	0
dense_2 (Dense)	(None, 64)	2112
dropout_1 (Dropout)	(None, 64)	0
dense_3 (Dense)	(None, 20)	1300

```
=====
==
```

```
Total params: 1,645,364
```

```
Trainable params: 1,645,364
```

```
Non-trainable params: 0
```

```
-----
--
```

```
None
```

```
In [77]:
```

```
r1 = model.fit(train_data_padded, train_labels, epochs=100,
batch_size=32, validation_data=(test_data_padded, test_labels),
class_weight=dict(inverse_class_weights))
```

```
Epoch 1/100
```

```
531/531 [=====] - 11s 20ms/step -
loss: 59.9142 - accuracy: 0.0312 - val_loss: 2.9933 -
val_accuracy: 0.0554
```

```
Epoch 2/100
```

```
531/531 [=====] - 4s 7ms/step - loss:
59.7056 - accuracy: 0.0884 - val_loss: 2.9779 - val_accuracy:
0.1603
```

Epoch 3/100

531/531 [=====] - 4s 7ms/step - loss: 59.3280 - accuracy: 0.0979 - val_loss: 2.9658 - val_accuracy: 0.1438

Epoch 4/100

531/531 [=====] - 3s 6ms/step - loss: 58.3765 - accuracy: 0.1104 - val_loss: 2.9028 - val_accuracy: 0.1219

Epoch 5/100

531/531 [=====] - 3s 5ms/step - loss: 56.0483 - accuracy: 0.1399 - val_loss: 2.7866 - val_accuracy: 0.1231

Epoch 6/100

531/531 [=====] - 3s 5ms/step - loss: 53.0405 - accuracy: 0.1491 - val_loss: 2.6679 - val_accuracy: 0.1681

Epoch 7/100

531/531 [=====] - 3s 5ms/step - loss: 49.8193 - accuracy: 0.1674 - val_loss: 2.5738 - val_accuracy: 0.1953

Epoch 8/100

531/531 [=====] - 4s 7ms/step - loss: 46.3996 - accuracy: 0.1842 - val_loss: 2.4852 - val_accuracy: 0.1870

Epoch 9/100

531/531 [=====] - 3s 5ms/step - loss:
43.9879 - accuracy: 0.1951 - val_loss: 2.4221 - val_accuracy:
0.2154

Epoch 10/100

531/531 [=====] - 3s 5ms/step - loss:
42.0127 - accuracy: 0.2095 - val_loss: 2.3389 - val_accuracy:
0.2351

Epoch 11/100

531/531 [=====] - 3s 5ms/step - loss:
40.6706 - accuracy: 0.2182 - val_loss: 2.2905 - val_accuracy:
0.2577

Epoch 12/100

531/531 [=====] - 3s 5ms/step - loss:
38.9891 - accuracy: 0.2344 - val_loss: 2.2401 - val_accuracy:
0.2723

Epoch 13/100

531/531 [=====] - 3s 5ms/step - loss:
37.4722 - accuracy: 0.2553 - val_loss: 2.1870 - val_accuracy:
0.2781

Epoch 14/100

531/531 [=====] - 3s 6ms/step - loss:
35.8749 - accuracy: 0.2728 - val_loss: 2.0974 - val_accuracy:
0.3082

Epoch 15/100

531/531 [=====] - 3s 5ms/step - loss:

33.9444 - accuracy: 0.3010 - val_loss: 2.0349 - val_accuracy:
0.3442

Epoch 16/100

531/531 [=====] - 3s 5ms/step - loss:
31.9808 - accuracy: 0.3286 - val_loss: 1.9389 - val_accuracy:
0.3777

Epoch 17/100

531/531 [=====] - 3s 5ms/step - loss:
29.8246 - accuracy: 0.3577 - val_loss: 1.8335 - val_accuracy:
0.4015

Epoch 18/100

531/531 [=====] - 3s 6ms/step - loss:
27.7647 - accuracy: 0.3832 - val_loss: 1.7307 - val_accuracy:
0.4282

Epoch 19/100

531/531 [=====] - 3s 6ms/step - loss:
25.5822 - accuracy: 0.4215 - val_loss: 1.6495 - val_accuracy:
0.4608

Epoch 20/100

531/531 [=====] - 3s 5ms/step - loss:
23.6537 - accuracy: 0.4559 - val_loss: 1.5950 - val_accuracy:
0.4892

Epoch 21/100

531/531 [=====] - 3s 5ms/step - loss:
21.8214 - accuracy: 0.4775 - val_loss: 1.5225 - val_accuracy:

0.5179

Epoch 22/100

531/531 [=====] - 3s 6ms/step - loss:
20.4309 - accuracy: 0.5091 - val_loss: 1.4603 - val_accuracy:
0.5239

Epoch 23/100

531/531 [=====] - 3s 5ms/step - loss:
18.7395 - accuracy: 0.5311 - val_loss: 1.4067 - val_accuracy:
0.5397

Epoch 24/100

531/531 [=====] - 3s 5ms/step - loss:
17.5805 - accuracy: 0.5568 - val_loss: 1.3477 - val_accuracy:
0.5669

Epoch 25/100

531/531 [=====] - 3s 5ms/step - loss:
16.2075 - accuracy: 0.5777 - val_loss: 1.2764 - val_accuracy:
0.6036

Epoch 26/100

531/531 [=====] - 3s 5ms/step - loss:
15.2806 - accuracy: 0.6019 - val_loss: 1.2530 - val_accuracy:
0.6148

Epoch 27/100

531/531 [=====] - 2s 5ms/step - loss:
14.0660 - accuracy: 0.6257 - val_loss: 1.1998 - val_accuracy:
0.6369

Epoch 28/100

531/531 [=====] - 3s 5ms/step - loss: 13.0900 - accuracy: 0.6486 - val_loss: 1.1602 - val_accuracy: 0.6476

Epoch 29/100

531/531 [=====] - 3s 5ms/step - loss: 12.3568 - accuracy: 0.6676 - val_loss: 1.1411 - val_accuracy: 0.6621

Epoch 30/100

531/531 [=====] - 3s 5ms/step - loss: 11.3669 - accuracy: 0.6905 - val_loss: 1.1007 - val_accuracy: 0.6680

Epoch 31/100

531/531 [=====] - 3s 6ms/step - loss: 10.6286 - accuracy: 0.7146 - val_loss: 1.0702 - val_accuracy: 0.6779

Epoch 32/100

531/531 [=====] - 3s 5ms/step - loss: 9.8552 - accuracy: 0.7281 - val_loss: 1.0374 - val_accuracy: 0.6944

Epoch 33/100

531/531 [=====] - 3s 5ms/step - loss: 9.3955 - accuracy: 0.7398 - val_loss: 1.0416 - val_accuracy: 0.6906

Epoch 34/100

531/531 [=====] - 3s 5ms/step - loss:
8.8528 - accuracy: 0.7568 - val_loss: 1.0104 - val_accuracy:
0.7066

Epoch 35/100

531/531 [=====] - 3s 5ms/step - loss:
8.2548 - accuracy: 0.7676 - val_loss: 0.9962 - val_accuracy:
0.7095

Epoch 36/100

531/531 [=====] - 3s 5ms/step - loss:
7.7419 - accuracy: 0.7846 - val_loss: 0.9886 - val_accuracy:
0.7127

Epoch 37/100

531/531 [=====] - 2s 5ms/step - loss:
7.3814 - accuracy: 0.7893 - val_loss: 0.9607 - val_accuracy:
0.7233

Epoch 38/100

531/531 [=====] - 3s 5ms/step - loss:
6.9352 - accuracy: 0.8069 - val_loss: 0.9699 - val_accuracy:
0.7280

Epoch 39/100

531/531 [=====] - 3s 5ms/step - loss:
6.3897 - accuracy: 0.8205 - val_loss: 0.9449 - val_accuracy:
0.7323

Epoch 40/100

531/531 [=====] - 3s 5ms/step - loss:

6.0012 - accuracy: 0.8366 - val_loss: 0.9583 - val_accuracy:
0.7294

Epoch 41/100

531/531 [=====] - 2s 5ms/step - loss:
5.6445 - accuracy: 0.8424 - val_loss: 0.9268 - val_accuracy:
0.7377

Epoch 42/100

531/531 [=====] - 3s 5ms/step - loss:
5.3952 - accuracy: 0.8474 - val_loss: 0.9267 - val_accuracy:
0.7457

Epoch 43/100

531/531 [=====] - 3s 7ms/step - loss:
5.2183 - accuracy: 0.8563 - val_loss: 0.9231 - val_accuracy:
0.7433

Epoch 44/100

531/531 [=====] - 3s 5ms/step - loss:
4.7608 - accuracy: 0.8657 - val_loss: 0.9196 - val_accuracy:
0.7491

Epoch 45/100

531/531 [=====] - 3s 5ms/step - loss:
4.4999 - accuracy: 0.8742 - val_loss: 0.9278 - val_accuracy:
0.7508

Epoch 46/100

531/531 [=====] - 2s 5ms/step - loss:
4.3075 - accuracy: 0.8802 - val_loss: 0.9163 - val_accuracy:

0.7522

Epoch 47/100

531/531 [=====] - 3s 5ms/step - loss:
4.3098 - accuracy: 0.8828 - val_loss: 0.9159 - val_accuracy:
0.7583

Epoch 48/100

531/531 [=====] - 3s 5ms/step - loss:
3.9671 - accuracy: 0.8904 - val_loss: 0.9177 - val_accuracy:
0.7549

Epoch 49/100

531/531 [=====] - 3s 5ms/step - loss:
3.7508 - accuracy: 0.8999 - val_loss: 0.9170 - val_accuracy:
0.7578

Epoch 50/100

531/531 [=====] - 3s 5ms/step - loss:
3.4167 - accuracy: 0.9050 - val_loss: 0.9120 - val_accuracy:
0.7654

Epoch 51/100

531/531 [=====] - 3s 5ms/step - loss:
3.3205 - accuracy: 0.9131 - val_loss: 0.9148 - val_accuracy:
0.7656

Epoch 52/100

531/531 [=====] - 2s 5ms/step - loss:
3.2107 - accuracy: 0.9130 - val_loss: 0.9166 - val_accuracy:
0.7675

Epoch 53/100

531/531 [=====] - 3s 5ms/step - loss: 3.0907 - accuracy: 0.9195 - val_loss: 0.9169 - val_accuracy: 0.7700

Epoch 54/100

531/531 [=====] - 2s 5ms/step - loss: 2.8278 - accuracy: 0.9243 - val_loss: 0.9265 - val_accuracy: 0.7685

Epoch 55/100

531/531 [=====] - 3s 6ms/step - loss: 2.7340 - accuracy: 0.9261 - val_loss: 0.9192 - val_accuracy: 0.7714

Epoch 56/100

531/531 [=====] - 2s 5ms/step - loss: 2.6732 - accuracy: 0.9307 - val_loss: 0.9404 - val_accuracy: 0.7724

Epoch 57/100

531/531 [=====] - 2s 5ms/step - loss: 2.5135 - accuracy: 0.9361 - val_loss: 0.9273 - val_accuracy: 0.7727

Epoch 58/100

531/531 [=====] - 3s 5ms/step - loss: 2.3983 - accuracy: 0.9387 - val_loss: 0.9354 - val_accuracy: 0.7722

Epoch 59/100

531/531 [=====] - 3s 5ms/step - loss:
2.2742 - accuracy: 0.9404 - val_loss: 0.9333 - val_accuracy:
0.7758

Epoch 60/100

531/531 [=====] - 3s 5ms/step - loss:
2.1182 - accuracy: 0.9481 - val_loss: 0.9434 - val_accuracy:
0.7795

Epoch 61/100

531/531 [=====] - 3s 5ms/step - loss:
2.0637 - accuracy: 0.9483 - val_loss: 0.9443 - val_accuracy:
0.7787

Epoch 62/100

531/531 [=====] - 3s 5ms/step - loss:
1.9177 - accuracy: 0.9512 - val_loss: 0.9614 - val_accuracy:
0.7780

Epoch 63/100

531/531 [=====] - 2s 5ms/step - loss:
1.9059 - accuracy: 0.9514 - val_loss: 0.9559 - val_accuracy:
0.7799

Epoch 64/100

531/531 [=====] - 3s 6ms/step - loss:
1.7550 - accuracy: 0.9569 - val_loss: 0.9626 - val_accuracy:
0.7833

Epoch 65/100

531/531 [=====] - 3s 6ms/step - loss:

1.7220 - accuracy: 0.9592 - val_loss: 0.9832 - val_accuracy:
0.7846

Epoch 66/100

531/531 [=====] - 3s 5ms/step - loss:
1.7179 - accuracy: 0.9596 - val_loss: 0.9819 - val_accuracy:
0.7833

Epoch 67/100

531/531 [=====] - 3s 6ms/step - loss:
1.5851 - accuracy: 0.9616 - val_loss: 0.9849 - val_accuracy:
0.7833

Epoch 68/100

531/531 [=====] - 3s 5ms/step - loss:
1.4692 - accuracy: 0.9647 - val_loss: 0.9851 - val_accuracy:
0.7829

Epoch 69/100

531/531 [=====] - 2s 5ms/step - loss:
1.5336 - accuracy: 0.9640 - val_loss: 0.9978 - val_accuracy:
0.7841

Epoch 70/100

531/531 [=====] - 2s 5ms/step - loss:
1.7107 - accuracy: 0.9606 - val_loss: 0.9974 - val_accuracy:
0.7831

Epoch 71/100

531/531 [=====] - 3s 5ms/step - loss:
1.5381 - accuracy: 0.9666 - val_loss: 1.0488 - val_accuracy:

0.7785

Epoch 72/100

531/531 [=====] - 3s 5ms/step - loss:
1.4126 - accuracy: 0.9662 - val_loss: 1.0092 - val_accuracy:
0.7858

Epoch 73/100

531/531 [=====] - 3s 5ms/step - loss:
1.3388 - accuracy: 0.9694 - val_loss: 1.0167 - val_accuracy:
0.7848

Epoch 74/100

531/531 [=====] - 3s 5ms/step - loss:
1.2163 - accuracy: 0.9700 - val_loss: 1.0332 - val_accuracy:
0.7848

Epoch 75/100

531/531 [=====] - 2s 5ms/step - loss:
1.1791 - accuracy: 0.9726 - val_loss: 1.0386 - val_accuracy:
0.7863

Epoch 76/100

531/531 [=====] - 3s 5ms/step - loss:
1.2157 - accuracy: 0.9727 - val_loss: 1.0261 - val_accuracy:
0.7875

Epoch 77/100

531/531 [=====] - 3s 5ms/step - loss:
1.2254 - accuracy: 0.9736 - val_loss: 1.0459 - val_accuracy:
0.7858

Epoch 78/100

531/531 [=====] - 3s 5ms/step - loss:
1.2307 - accuracy: 0.9706 - val_loss: 1.0429 - val_accuracy:
0.7858

Epoch 79/100

531/531 [=====] - 3s 6ms/step - loss:
1.2113 - accuracy: 0.9738 - val_loss: 1.0581 - val_accuracy:
0.7863

Epoch 80/100

531/531 [=====] - 3s 5ms/step - loss:
1.0248 - accuracy: 0.9775 - val_loss: 1.0508 - val_accuracy:
0.7872

Epoch 81/100

531/531 [=====] - 3s 5ms/step - loss:
1.0278 - accuracy: 0.9756 - val_loss: 1.0654 - val_accuracy:
0.7884

Epoch 82/100

531/531 [=====] - 2s 5ms/step - loss:
0.9587 - accuracy: 0.9783 - val_loss: 1.0769 - val_accuracy:
0.7901

Epoch 83/100

531/531 [=====] - 2s 5ms/step - loss:
0.9815 - accuracy: 0.9786 - val_loss: 1.0817 - val_accuracy:
0.7867

Epoch 84/100

531/531 [=====] - 3s 5ms/step - loss:
0.9486 - accuracy: 0.9788 - val_loss: 1.0876 - val_accuracy:
0.7884

Epoch 85/100

531/531 [=====] - 3s 5ms/step - loss:
0.8395 - accuracy: 0.9823 - val_loss: 1.0883 - val_accuracy:
0.7880

Epoch 86/100

531/531 [=====] - 3s 5ms/step - loss:
0.8297 - accuracy: 0.9818 - val_loss: 1.1006 - val_accuracy:
0.7853

Epoch 87/100

531/531 [=====] - 2s 5ms/step - loss:
0.8609 - accuracy: 0.9808 - val_loss: 1.0987 - val_accuracy:
0.7860

Epoch 88/100

531/531 [=====] - 3s 5ms/step - loss:
0.8015 - accuracy: 0.9838 - val_loss: 1.1071 - val_accuracy:
0.7875

Epoch 89/100

531/531 [=====] - 3s 5ms/step - loss:
0.7211 - accuracy: 0.9834 - val_loss: 1.1173 - val_accuracy:
0.7892

Epoch 90/100

531/531 [=====] - 2s 5ms/step - loss:

0.8440 - accuracy: 0.9818 - val_loss: 1.1224 - val_accuracy:
0.7892

Epoch 91/100

531/531 [=====] - 3s 6ms/step - loss:
0.7433 - accuracy: 0.9836 - val_loss: 1.1264 - val_accuracy:
0.7870

Epoch 92/100

531/531 [=====] - 2s 5ms/step - loss:
0.7630 - accuracy: 0.9842 - val_loss: 1.1179 - val_accuracy:
0.7880

Epoch 93/100

531/531 [=====] - 3s 5ms/step - loss:
0.6864 - accuracy: 0.9861 - val_loss: 1.1443 - val_accuracy:
0.7875

Epoch 94/100

531/531 [=====] - 2s 5ms/step - loss:
0.7292 - accuracy: 0.9845 - val_loss: 1.1516 - val_accuracy:
0.7865

Epoch 95/100

531/531 [=====] - 3s 5ms/step - loss:
0.6997 - accuracy: 0.9850 - val_loss: 1.1412 - val_accuracy:
0.7870

Epoch 96/100

531/531 [=====] - 2s 5ms/step - loss:
0.6978 - accuracy: 0.9852 - val_loss: 1.1373 - val_accuracy:

0.7889

Epoch 97/100

531/531 [=====] - 3s 5ms/step - loss:
0.6556 - accuracy: 0.9852 - val_loss: 1.1587 - val_accuracy:
0.7858

Epoch 98/100

531/531 [=====] - 2s 5ms/step - loss:
0.6145 - accuracy: 0.9880 - val_loss: 1.1610 - val_accuracy:
0.7894

Epoch 99/100

531/531 [=====] - 2s 5ms/step - loss:
0.6179 - accuracy: 0.9866 - val_loss: 1.1715 - val_accuracy:
0.7860

Epoch 100/100


531/531 [=====] - 3s 5ms/step - loss:
0.6324 - accuracy: 0.9862 - val_loss: 1.1672 - val_accuracy:
0.7870

In [78]:

train also a simple ANN

model2 = make_model_ann()

print(model2.summary())



Model: "model_2"

--		
Layer (type)	Output Shape	Param #
=====		
==		
input_3 (InputLayer)	[(None, 18000)]	0
dense_4 (Dense)	(None, 1028)	18505028
dropout_2 (Dropout)	(None, 1028)	0
dense_5 (Dense)	(None, 512)	526848
dropout_3 (Dropout)	(None, 512)	0
dense_6 (Dense)	(None, 256)	131328
dense_7 (Dense)	(None, 20)	5140
=====		
==		
Total params: 19,168,344		
Trainable params: 19,168,344		
Non-trainable params: 0		

--
None

In [79]:

```
r2 = model2.fit(train_data_ann.toarray(), train_labels,  
epochs=100, batch_size=32,  
validation_data=(test_data_ann.toarray(), test_labels),  
class_weight=dict(inverse_class_weights), shuffle=True)
```

Epoch 1/100

531/531 [=====] - 7s 11ms/step - loss: 59.3496 - accuracy: 0.1099 - val_loss: 2.9208 - val_accuracy: 0.2652

Epoch 2/100

531/531 [=====] - 5s 8ms/step - loss: 53.0952 - accuracy: 0.2433 - val_loss: 2.4742 - val_accuracy: 0.3384

Epoch 3/100

531/531 [=====] - 4s 8ms/step - loss: 33.7288 - accuracy: 0.4246 - val_loss: 1.6334 - val_accuracy: 0.5373

Epoch 4/100

531/531 [=====] - 4s 8ms/step - loss:
19.9803 - accuracy: 0.5876 - val_loss: 1.2084 - val_accuracy:
0.6738

Epoch 5/100

531/531 [=====] - 4s 8ms/step - loss:
13.7646 - accuracy: 0.6871 - val_loss: 1.0191 - val_accuracy:
0.7037

Epoch 6/100

531/531 [=====] - 4s 8ms/step - loss:
10.4509 - accuracy: 0.7469 - val_loss: 0.8996 - val_accuracy:
0.7345

Epoch 7/100

531/531 [=====] - 5s 9ms/step - loss:
8.1547 - accuracy: 0.7928 - val_loss: 0.8309 - val_accuracy:
0.7561

Epoch 8/100

531/531 [=====] - 4s 8ms/step - loss:
6.6510 - accuracy: 0.8230 - val_loss: 0.7322 - val_accuracy:
0.7802

Epoch 9/100

531/531 [=====] - 4s 8ms/step - loss:
5.4875 - accuracy: 0.8501 - val_loss: 0.6871 - val_accuracy:
0.7928

Epoch 10/100

531/531 [=====] - 4s 8ms/step - loss:

4.5912 - accuracy: 0.8733 - val_loss: 0.6747 - val_accuracy:
0.7962

Epoch 11/100

531/531 [=====] - 4s 8ms/step - loss:
3.8537 - accuracy: 0.8879 - val_loss: 0.6750 - val_accuracy:
0.8006

Epoch 12/100

531/531 [=====] - 5s 9ms/step - loss:
3.2754 - accuracy: 0.9083 - val_loss: 0.6381 - val_accuracy:
0.8115

Epoch 13/100

531/531 [=====] - 4s 8ms/step - loss:
2.8264 - accuracy: 0.9181 - val_loss: 0.6324 - val_accuracy:
0.8190

Epoch 14/100

531/531 [=====] - 5s 9ms/step - loss:
2.4749 - accuracy: 0.9306 - val_loss: 0.6374 - val_accuracy:
0.8220

Epoch 15/100

531/531 [=====] - 5s 9ms/step - loss:
2.1491 - accuracy: 0.9381 - val_loss: 0.6302 - val_accuracy:
0.8229

Epoch 16/100

531/531 [=====] - 4s 8ms/step - loss:
1.9063 - accuracy: 0.9468 - val_loss: 0.6319 - val_accuracy:

0.8278

Epoch 17/100

531/531 [=====] - 4s 8ms/step - loss:
1.6878 - accuracy: 0.9530 - val_loss: 0.6385 - val_accuracy:
0.8283

Epoch 18/100

531/531 [=====] - 5s 9ms/step - loss:
1.5010 - accuracy: 0.9592 - val_loss: 0.6416 - val_accuracy:
0.8297

Epoch 19/100

531/531 [=====] - 4s 8ms/step - loss:
1.3219 - accuracy: 0.9650 - val_loss: 0.6553 - val_accuracy:
0.8285

Epoch 20/100

531/531 [=====] - 4s 8ms/step - loss:
1.1601 - accuracy: 0.9689 - val_loss: 0.6488 - val_accuracy:
0.8317

Epoch 21/100

531/531 [=====] - 4s 8ms/step - loss:
1.0776 - accuracy: 0.9712 - val_loss: 0.6625 - val_accuracy:
0.8271

Epoch 22/100

531/531 [=====] - 5s 9ms/step - loss:
0.9872 - accuracy: 0.9752 - val_loss: 0.6662 - val_accuracy:
0.8302

Epoch 23/100

531/531 [=====] - 4s 8ms/step - loss:
0.8918 - accuracy: 0.9770 - val_loss: 0.6763 - val_accuracy:
0.8288

Epoch 24/100

531/531 [=====] - 5s 8ms/step - loss:
0.8101 - accuracy: 0.9812 - val_loss: 0.6719 - val_accuracy:
0.8314

Epoch 25/100

531/531 [=====] - 4s 8ms/step - loss:
0.7049 - accuracy: 0.9839 - val_loss: 0.6813 - val_accuracy:
0.8295

Epoch 26/100

531/531 [=====] - 4s 8ms/step - loss:
0.6768 - accuracy: 0.9841 - val_loss: 0.6940 - val_accuracy:
0.8297

Epoch 27/100

531/531 [=====] - 4s 8ms/step - loss:
0.6347 - accuracy: 0.9855 - val_loss: 0.7071 - val_accuracy:
0.8297

Epoch 28/100

531/531 [=====] - 4s 8ms/step - loss:
0.5714 - accuracy: 0.9881 - val_loss: 0.7053 - val_accuracy:
0.8297

Epoch 29/100

531/531 [=====] - 5s 9ms/step - loss:
0.5252 - accuracy: 0.9892 - val_loss: 0.7156 - val_accuracy:
0.8297

Epoch 30/100

531/531 [=====] - 4s 8ms/step - loss:
0.5076 - accuracy: 0.9895 - val_loss: 0.7211 - val_accuracy:
0.8297

Epoch 31/100

531/531 [=====] - 4s 8ms/step - loss:
0.4791 - accuracy: 0.9906 - val_loss: 0.7302 - val_accuracy:
0.8317

Epoch 32/100

531/531 [=====] - 4s 8ms/step - loss:
0.4032 - accuracy: 0.9917 - val_loss: 0.7361 - val_accuracy:
0.8317

Epoch 33/100

531/531 [=====] - 4s 8ms/step - loss:
0.4039 - accuracy: 0.9919 - val_loss: 0.7441 - val_accuracy:
0.8261

Epoch 34/100

531/531 [=====] - 4s 8ms/step - loss:
0.3883 - accuracy: 0.9923 - val_loss: 0.7457 - val_accuracy:
0.8292

Epoch 35/100

531/531 [=====] - 4s 8ms/step - loss:

0.3482 - accuracy: 0.9942 - val_loss: 0.7600 - val_accuracy:
0.8275

Epoch 36/100

531/531 [=====] - 5s 9ms/step - loss:
0.3278 - accuracy: 0.9940 - val_loss: 0.7609 - val_accuracy:
0.8295

Epoch 37/100

531/531 [=====] - 4s 8ms/step - loss:
0.3191 - accuracy: 0.9942 - val_loss: 0.7683 - val_accuracy:
0.8290

Epoch 38/100

531/531 [=====] - 4s 8ms/step - loss:
0.3362 - accuracy: 0.9943 - val_loss: 0.7717 - val_accuracy:
0.8288

Epoch 39/100

531/531 [=====] - 4s 8ms/step - loss:
0.2956 - accuracy: 0.9947 - val_loss: 0.7717 - val_accuracy:
0.8302

Epoch 40/100

531/531 [=====] - 4s 8ms/step - loss:
0.3042 - accuracy: 0.9942 - val_loss: 0.7848 - val_accuracy:
0.8326

Epoch 41/100

531/531 [=====] - 4s 8ms/step - loss:
0.2688 - accuracy: 0.9954 - val_loss: 0.7873 - val_accuracy:

0.8317

Epoch 42/100

531/531 [=====] - 4s 8ms/step - loss:
0.2625 - accuracy: 0.9953 - val_loss: 0.7924 - val_accuracy:
0.8302

Epoch 43/100

531/531 [=====] - 5s 8ms/step - loss:
0.2561 - accuracy: 0.9954 - val_loss: 0.7943 - val_accuracy:
0.8297

Epoch 44/100

531/531 [=====] - 5s 9ms/step - loss:
0.2245 - accuracy: 0.9961 - val_loss: 0.7995 - val_accuracy:
0.8317

Epoch 45/100

531/531 [=====] - 4s 8ms/step - loss:
0.2569 - accuracy: 0.9954 - val_loss: 0.8092 - val_accuracy:
0.8292

Epoch 46/100

531/531 [=====] - 4s 8ms/step - loss:
0.2410 - accuracy: 0.9959 - val_loss: 0.8092 - val_accuracy:
0.8314

Epoch 47/100

531/531 [=====] - 4s 8ms/step - loss:
0.2194 - accuracy: 0.9959 - val_loss: 0.8133 - val_accuracy:
0.8275

Epoch 48/100

531/531 [=====] - 4s 8ms/step - loss:
0.2042 - accuracy: 0.9966 - val_loss: 0.8152 - val_accuracy:
0.8297

Epoch 49/100

531/531 [=====] - 5s 9ms/step - loss:
0.1949 - accuracy: 0.9964 - val_loss: 0.8224 - val_accuracy:
0.8288

Epoch 50/100

531/531 [=====] - 5s 9ms/step - loss:
0.2053 - accuracy: 0.9964 - val_loss: 0.8172 - val_accuracy:
0.8300

Epoch 51/100

531/531 [=====] - 4s 8ms/step - loss:
0.2076 - accuracy: 0.9961 - val_loss: 0.8298 - val_accuracy:
0.8307

Epoch 52/100

531/531 [=====] - 4s 8ms/step - loss:
0.1786 - accuracy: 0.9970 - val_loss: 0.8308 - val_accuracy:
0.8309

Epoch 53/100

531/531 [=====] - 4s 8ms/step - loss:
0.1694 - accuracy: 0.9974 - val_loss: 0.8380 - val_accuracy:
0.8319

Epoch 54/100

531/531 [=====] - 5s 9ms/step - loss:
0.1737 - accuracy: 0.9966 - val_loss: 0.8390 - val_accuracy:
0.8300

Epoch 55/100

531/531 [=====] - 4s 8ms/step - loss:
0.1752 - accuracy: 0.9969 - val_loss: 0.8389 - val_accuracy:
0.8271

Epoch 56/100

531/531 [=====] - 4s 8ms/step - loss:
0.1638 - accuracy: 0.9973 - val_loss: 0.8438 - val_accuracy:
0.8307

Epoch 57/100

531/531 [=====] - 5s 9ms/step - loss:
0.1627 - accuracy: 0.9970 - val_loss: 0.8474 - val_accuracy:
0.8309

Epoch 58/100

531/531 [=====] - 4s 8ms/step - loss:
0.1621 - accuracy: 0.9972 - val_loss: 0.8518 - val_accuracy:
0.8329

Epoch 59/100

531/531 [=====] - 5s 9ms/step - loss:
0.1530 - accuracy: 0.9974 - val_loss: 0.8524 - val_accuracy:
0.8319

Epoch 60/100

531/531 [=====] - 4s 8ms/step - loss:

0.1619 - accuracy: 0.9968 - val_loss: 0.8556 - val_accuracy:
0.8317

Epoch 61/100

531/531 [=====] - 4s 8ms/step - loss:
0.1450 - accuracy: 0.9976 - val_loss: 0.8704 - val_accuracy:
0.8300

Epoch 62/100

531/531 [=====] - 4s 8ms/step - loss:
0.1656 - accuracy: 0.9966 - val_loss: 0.8672 - val_accuracy:
0.8275

Epoch 63/100

531/531 [=====] - 4s 8ms/step - loss:
0.1386 - accuracy: 0.9976 - val_loss: 0.8645 - val_accuracy:
0.8324

Epoch 64/100

531/531 [=====] - 5s 9ms/step - loss:
0.1377 - accuracy: 0.9975 - val_loss: 0.8697 - val_accuracy:
0.8285

Epoch 65/100

531/531 [=====] - 5s 8ms/step - loss:
0.1511 - accuracy: 0.9972 - val_loss: 0.8693 - val_accuracy:
0.8331

Epoch 66/100

531/531 [=====] - 4s 8ms/step - loss:
0.1450 - accuracy: 0.9975 - val_loss: 0.8744 - val_accuracy:

0.8314

Epoch 67/100

531/531 [=====] - 4s 8ms/step - loss:
0.1413 - accuracy: 0.9976 - val_loss: 0.8756 - val_accuracy:
0.8302

Epoch 68/100

531/531 [=====] - 4s 8ms/step - loss:
0.1326 - accuracy: 0.9974 - val_loss: 0.8799 - val_accuracy:
0.8273

Epoch 69/100

531/531 [=====] - 5s 9ms/step - loss:
0.1425 - accuracy: 0.9975 - val_loss: 0.8755 - val_accuracy:
0.8317

Epoch 70/100

531/531 [=====] - 4s 8ms/step - loss:
0.1448 - accuracy: 0.9972 - val_loss: 0.8793 - val_accuracy:
0.8305

Epoch 71/100

531/531 [=====] - 4s 8ms/step - loss:
0.1185 - accuracy: 0.9975 - val_loss: 0.8894 - val_accuracy:
0.8292

Epoch 72/100

531/531 [=====] - 5s 10ms/step - loss:
0.1233 - accuracy: 0.9979 - val_loss: 0.8914 - val_accuracy:
0.8288

Epoch 73/100

531/531 [=====] - 4s 8ms/step - loss:
0.1395 - accuracy: 0.9978 - val_loss: 0.8879 - val_accuracy:
0.8319

Epoch 74/100

531/531 [=====] - 4s 8ms/step - loss:
0.1126 - accuracy: 0.9981 - val_loss: 0.8919 - val_accuracy:
0.8300

Epoch 75/100

531/531 [=====] - 4s 8ms/step - loss:
0.1303 - accuracy: 0.9977 - val_loss: 0.8992 - val_accuracy:
0.8324

Epoch 76/100

531/531 [=====] - 4s 8ms/step - loss:
0.1247 - accuracy: 0.9975 - val_loss: 0.9046 - val_accuracy:
0.8288

Epoch 77/100

531/531 [=====] - 4s 8ms/step - loss:
0.1217 - accuracy: 0.9978 - val_loss: 0.8998 - val_accuracy:
0.8297

Epoch 78/100

531/531 [=====] - 4s 8ms/step - loss:
0.1256 - accuracy: 0.9976 - val_loss: 0.9039 - val_accuracy:
0.8319

Epoch 79/100

531/531 [=====] - 5s 9ms/step - loss:
0.1284 - accuracy: 0.9974 - val_loss: 0.9012 - val_accuracy:
0.8305

Epoch 80/100

531/531 [=====] - 4s 8ms/step - loss:
0.1204 - accuracy: 0.9979 - val_loss: 0.9025 - val_accuracy:
0.8297

Epoch 81/100

531/531 [=====] - 4s 8ms/step - loss:
0.1121 - accuracy: 0.9980 - val_loss: 0.9131 - val_accuracy:
0.8275

Epoch 82/100

531/531 [=====] - 4s 8ms/step - loss:
0.1153 - accuracy: 0.9980 - val_loss: 0.9135 - val_accuracy:
0.8309

Epoch 83/100

531/531 [=====] - 4s 8ms/step - loss:
0.1092 - accuracy: 0.9981 - val_loss: 0.9152 - val_accuracy:
0.8312

Epoch 84/100

531/531 [=====] - 4s 8ms/step - loss:
0.0935 - accuracy: 0.9984 - val_loss: 0.9175 - val_accuracy:
0.8290

Epoch 85/100

531/531 [=====] - 4s 8ms/step - loss:

0.1042 - accuracy: 0.9979 - val_loss: 0.9181 - val_accuracy:
0.8295

Epoch 86/100

531/531 [=====] - 5s 9ms/step - loss:
0.1098 - accuracy: 0.9981 - val_loss: 0.9175 - val_accuracy:
0.8302

Epoch 87/100

531/531 [=====] - 5s 9ms/step - loss:
0.1022 - accuracy: 0.9981 - val_loss: 0.9198 - val_accuracy:
0.8319

Epoch 88/100

531/531 [=====] - 4s 8ms/step - loss:
0.1036 - accuracy: 0.9980 - val_loss: 0.9220 - val_accuracy:
0.8324

Epoch 89/100

531/531 [=====] - 4s 8ms/step - loss:
0.1143 - accuracy: 0.9977 - val_loss: 0.9240 - val_accuracy:
0.8307

Epoch 90/100

531/531 [=====] - 4s 8ms/step - loss:
0.1047 - accuracy: 0.9982 - val_loss: 0.9324 - val_accuracy:
0.8285

Epoch 91/100

531/531 [=====] - 4s 8ms/step - loss:
0.1002 - accuracy: 0.9982 - val_loss: 0.9309 - val_accuracy:

0.8268

Epoch 92/100

531/531 [=====] - 5s 9ms/step - loss:
0.1052 - accuracy: 0.9979 - val_loss: 0.9309 - val_accuracy:
0.8273

Epoch 93/100

531/531 [=====] - 4s 8ms/step - loss:
0.0919 - accuracy: 0.9979 - val_loss: 0.9342 - val_accuracy:
0.8317

Epoch 94/100

531/531 [=====] - 4s 8ms/step - loss:
0.0916 - accuracy: 0.9981 - val_loss: 0.9303 - val_accuracy:
0.8307

Epoch 95/100

531/531 [=====] - 4s 8ms/step - loss:
0.0977 - accuracy: 0.9982 - val_loss: 0.9397 - val_accuracy:
0.8275

Epoch 96/100

531/531 [=====] - 4s 8ms/step - loss:
0.0966 - accuracy: 0.9981 - val_loss: 0.9381 - val_accuracy:
0.8275

Epoch 97/100

531/531 [=====] - 4s 8ms/step - loss:
0.0846 - accuracy: 0.9984 - val_loss: 0.9384 - val_accuracy:
0.8283

Epoch 98/100

531/531 [=====] - 4s 8ms/step - loss: 0.0956 - accuracy: 0.9980 - val_loss: 0.9400 - val_accuracy: 0.8305

Epoch 99/100

531/531 [=====] - 4s 8ms/step - loss: 0.0883 - accuracy: 0.9982 - val_loss: 0.9500 - val_accuracy: 0.8278

Epoch 100/100

531/531 [=====] - 4s 8ms/step - loss: 0.0941 - accuracy: 0.9982 - val_loss: 0.9412 - val_accuracy: 0.8288

6. Evaluation

Note: A sequence model like RNN or LSTM was not considered because it was overfitting strongly

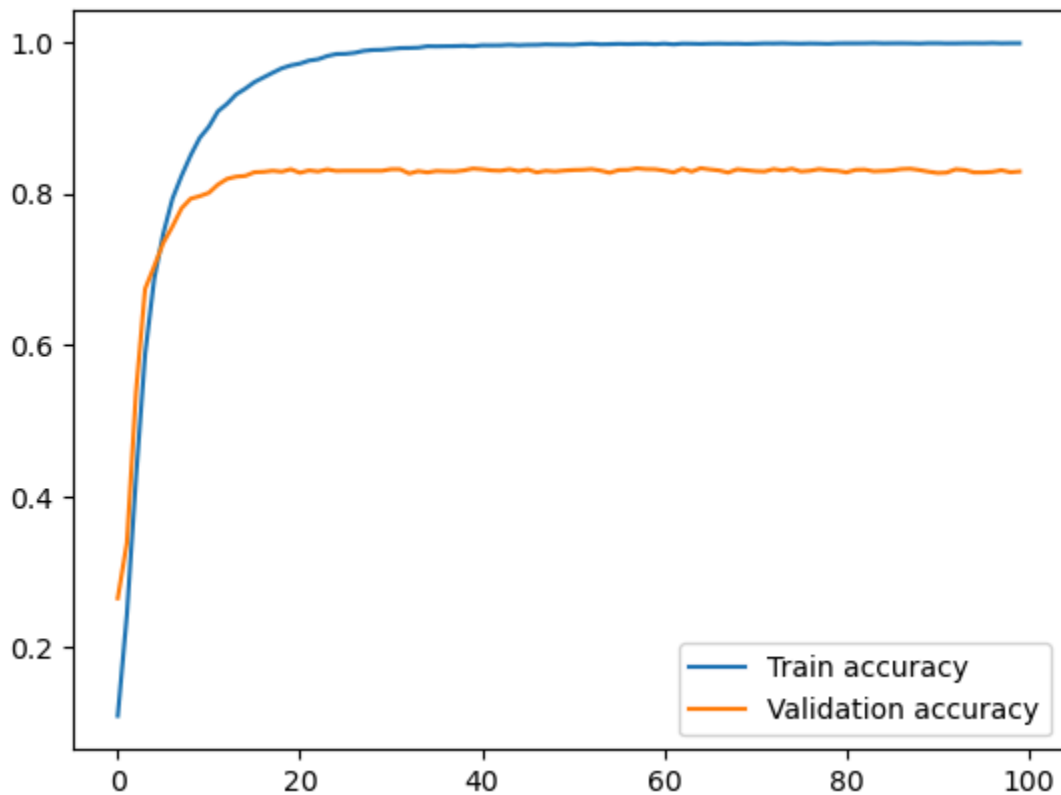
The CNN performed much better for the sequence and Embedding data

6.1 Plain Feedforward Network Evaluation

In [80]:


```
plt.plot(r2.history["accuracy"], label="Train accuracy")
plt.plot(r2.history["val_accuracy"], label="Validation accuracy")
plt.legend()
```

```
plt.show()
```



In [81]:

```
plt.plot(r2.history["loss"], label="Train loss")  
plt.plot(r2.history["val_loss"], label="Validation loss")  
plt.legend()  
plt.show()
```

In [82]:

```
from sklearn.metrics import f1_score, classification_report
preds = np.argmax(model2.predict(test_data_ann), axis=1)
```

```
129/129 [=====] - 0s 3ms/step
```

In [83]:

```
print("Classification report for test data")
print(classification_report(preds, test_labels))
```

Classification report for test data

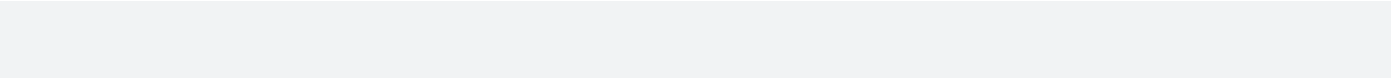
	precision	recall	f1-score	support
0	0.79	0.82	0.81	71
1	0.86	0.86	0.86	215
2	0.86	0.82	0.84	888
3	0.77	0.88	0.82	67
4	0.99	0.95	0.97	101
5	0.94	0.91	0.92	250
6	0.82	0.82	0.82	147
7	0.85	0.86	0.85	159
8	0.84	0.77	0.81	35
9	0.76	0.71	0.73	363
10	0.69	0.41	0.51	22
11	0.93	0.87	0.90	15
12	0.81	0.88	0.84	109
13	0.72	0.74	0.73	113
14	0.82	0.82	0.82	416
15	0.78	0.76	0.77	129
16	0.88	0.92	0.90	239
17	0.81	0.88	0.85	103
18	0.81	0.87	0.84	487
19	0.71	0.74	0.73	188
accuracy			0.83	4117


macro avg	0.82	0.81	0.82	4117
weighted avg	0.83	0.83	0.83	4117

6.2 Plain Feedforward Network Evaluation

In [84]:


```
plt.plot(r1.history["accuracy"], label="Train accuracy")
plt.plot(r1.history["val_accuracy"], label="Validation
accuracy")
plt.legend()
plt.show()
```





In [85]:

```
plt.plot(r1.history["loss"], label="Train loss")
plt.plot(r1.history["val_loss"], label="Validation loss")
plt.legend()
plt.show()
```



In [86]:

```
preds = np.argmax(model .predict(test_data_padded), axis=1)
print("Classification report for test data")
print(classification_report(preds,test_labels))
```

129/129 [=====] - 0s 2ms/step

Classification report for test data

	precision	recall	f1-score	support
0	0.41	0.88	0.56	34
1	0.79	0.87	0.82	194
2	0.84	0.74	0.79	962

3	0.74	0.79	0.77	72
4	0.96	0.97	0.96	96
5	0.93	0.89	0.91	254
6	0.79	0.80	0.79	144
7	0.82	0.82	0.82	161
8	0.88	0.76	0.81	37
9	0.72	0.68	0.70	360
10	0.77	0.48	0.59	21
11	0.93	0.81	0.87	16
12	0.73	0.75	0.74	116
13	0.62	0.88	0.73	82
14	0.78	0.76	0.77	425
15	0.75	0.79	0.77	119
16	0.83	0.84	0.84	246
17	0.76	0.93	0.84	91
18	0.76	0.78	0.77	515
19	0.72	0.83	0.77	172
accuracy			0.79	4117
macro avg	0.78	0.80	0.78	4117
weighted avg	0.79	0.79	0.79	4117

[Reference link](#)